Wouter. Your first week at Ditto. A warm welcome!

Name: Wouter van Wessel · **Role:** Growth Engineer · **Lead:** Niek Schreurs

---

Why you were hired

Growth (engineering) is not marketing, and it's not traditional product engineering. It sits at the intersection of product, data, and experimentation.

A product engineer builds features. A growth engineer asks whether those features actually change behavior — and proves it. You'll spend your time instrumenting events so we can measure what users do, designing experiments to test why they do it, analyzing whether an intervention actually moved the number or just looked like it did, and feeding those learnings back into the next cycle.

The mindset is scientific. You start with a hypothesis, not a feature idea. You define what success looks like before you ship, not after. You care about causation, not just correlation. And you think in systems — because changing one part of the funnel always affects the others.

At Ditto, the growth engineer works directly with Niek, Head of Growth, and partners closely with product and design on the features that matter most for getting users to value. You'll have a lot of autonomy and a lot of leverage. The tools, the data, and the AI systems are already in place. What's missing is someone who can run the experiments and build the growth model.

So, let's dive a bit deeper into Ditto and the numbers.

Ditto closed its seed round in Q2 2026. We have a product people want, meaning there is product-market fit — 4,349 installs every month, and 45% of them register. That's well above any benchmark. The demand is there. The problem is what happens after.

Of everyone who registers, only 12% view a summary within their first 14 days. Our Q2 target is to bring that to 25%. Customer.io is going to help a lot here. Viewing a summary is the moment Ditto clicks — the moment a conversation with your doctor turns into something you can actually understand, share with your partner, and act on. People who get there come back. People who view two summaries in their first month retain at dramatically higher rates. But right now, 88% of registered users never see that moment. They showed up with intent, and we lost them.

This isn't new, and it's not getting worse. It's actually getting better. Back in January, D14 activation was at 10% on the registration side and 4% from install. Since then, it's been climbing

— March peaked at 15% and 6%, with some days in late March hitting 20-28% on the registration funnel. April is holding around 14% and 7%. The trajectory is up, but we're starting from a low base, and the north star is far away.

Core MAU today (29 April) is 1,739, below our linear target of **2,165, -462** compared to today's pace (**-21.4%**) — people who viewed a summary in the last 30 days. The Q2 exit target is 4,568. The end-of-year target is 50,000. We're at 3.5% of where we need to be in eight months. We are not going to close that gap by pouring more people into a funnel that loses 88% (insights we will discuss) of them before they see value. The path to 50K runs through activation and retention first, acquisition second.

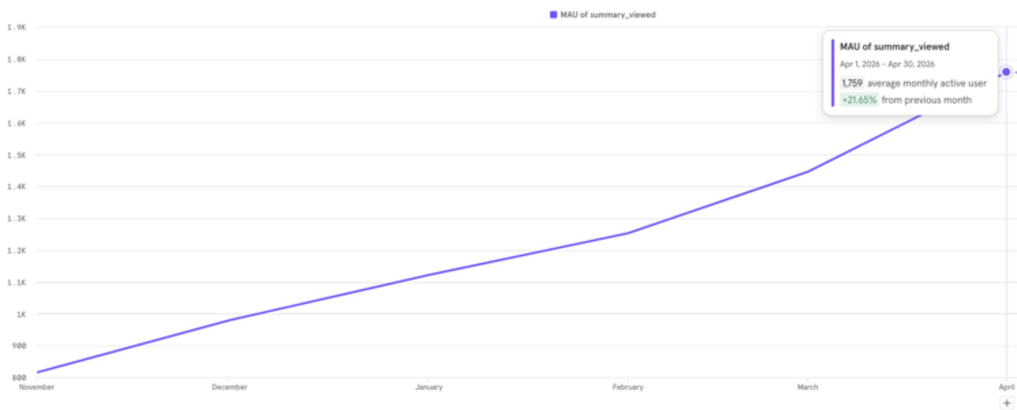

We also know where the biggest single leak is. 55% of users who add an appointment manually — who literally prepare for the moment that creates a summary — never view a summary within 60 days. They did the work but never got the payoff. That gap is the most obvious place to start. We pushed a few initiatives for that:

- V1 - Push to Calendar
- V2 - Calendar Auto Sync

V1 — Push to Calendar (launched ~Mar 23)

- Adoption: **5%** from register (14% see explanation screen → 33% add calendar event)
- Downstream: 87% create manual appt → **29% view a summary** within 14d
- That 29% is **2.2× the baseline** D14 rate of 13%
- Contribution: ~61 activated users in the Feb–Apr 15 window — roughly **6% of all activations**
- Calendar event adders are high-intent users — the feature is a quality filter, not a volume driver

V2 — Calendar Auto Sync (launched ~Apr 13–20)

- Adoption: **6%** from register (16% see explanation → 35% enable auto sync)
- Downstream: 96% create manual appt (higher than V1's 87%) → **13% view summary** so far
- The 13% is immature — V2 cohort is <2 weeks old and manual appt → summary has a bimodal lag (52% same-day, rest trickles over weeks). Expect this to climb to 20-25% at maturity
- 263 users enabled auto sync already — volume is promising

Net effect on KR2 (D14 activation 11% → 25%)

- Combined calendar features contribute ~20-25% of all manual appointment activations
- But overall adoption remains low (~5-6% from the register). The **60% skip rate on the explanation screen** is the biggest lever — halving that skip rate would nearly double calendar-driven activations
- The big D14 jump from ~10% (Jan) to ~18-20% (late Mar/Apr) was primarily driven by Care Circle v2 (Mar 17, +8pp on summary viewing) — calendar features were a contributor, not the main driver

Bottom line: Both features improve activation quality for users who adopt them, but at 5-6% adoption, they move the needle by ~1-2pp on D14. The unlock is reducing the skip rate on that explanation screen — that's where the volume is trapped.

But solving activation isn't just an individual journey. Ditto has a care circle — a shared space where family members and caretakers can follow someone's medical care. Within a care circle, you can view summaries created by someone else. **That means a caretaker who never records a conversation themselves can still hit the aha moment through their parent's or partner's appointment.** The care circle is both a retention mechanism and a second path to activation. How big that path is, and how to optimize it, is something we're still figuring out.

That's the honest picture. We have one strong signal — summary equals retention. We have a clear gap — 88% of registered users don't get there. We have an upward trend that shows the product is improving. And we have a set of hypotheses about care circles, referral loops, and retention mechanics that haven't been validated yet. The full growth model doesn't exist. Building it is the job.

You're here to run the experiments that grow Ditto. Design and ship A/B tests. Find the correlations that actually move the MAU needle. Separate signal from noise in the data. Together with Niek, you'll build the growth model — figuring out which behaviors predict retention, which interventions change those behaviors, and how the pieces connect into a system that compounds.

You're not joining a team that has the playbook figured out and needs someone to execute it. You're joining to help build the playbook — one experiment at a time.

Together with Niek, you will be driving the following two KRs for Q2:

Objective 1
Build the right team, habits, and foundation that unlock hyper scaling Ditto internationally.

KR4: 100% of significant initiatives follow a thorough Proposal & Evaluation.

Objective 2
Build the growth loop that compounds into 5K MAU by 20260630

KR2: D14 activation rate goes from 11% to 25%

As you likely understand, objective 2 is a core growth focus and, as noted, is behind target (although we grow 15-20% MoM - we are super ambitious)

Growth at Ditto

To get an understanding of how we "do" Growth at Ditto. Please read the following document. Comment with any questions to help us learn, clarify, and build a shared source of truth for growth together.

Your week-1 key result

By Friday, you should be able to answer two questions: what are we building and why, and how do we measure whether it's working.

Understand the product roadmap

Spend time with Merlijn (CPO) and Niek this week to understand what Ditto is building and where it's heading. This isn't a passive briefing — come with questions 🙂 . You'll want to understand which features are in flight right now, what problem each one solves, and how they connect to activation and retention (OKRs). Ask Merlijn about the product vision and the choices behind what's on the roadmap and what isn't. Ask Niek about the growth priorities and where the biggest gaps are in the funnel. By the end of the week, you should have a mental map of how product decisions and growth objectives connect — because that's the intersection you'll operate in every day.

Talk to Oli about international expansion

We're launching in the UK soon and Germany is on the horizon. Sit down with Ollie this week to understand the timeline, what's different about those markets, and what we know so far about user behavior outside The Netherlands.

This matters for your work directly. The experiments you design need to account for the fact that what works in the Dutch market may not transfer. Language, healthcare systems, cultural expectations around medical care and family involvement — all of these affect activation. A summary reminder flow that works in The Netherlands might fall flat in the UK if the care circle dynamic is different there.

Come out of that conversation with a sense of what questions expansion raises for experimentation. Do we need to validate our aha moment in new markets, or can we assume it holds? Should early experiments run per-market or globally? What would a UK-specific activation funnel look like — is it the same shape, or does the entry point change when the healthcare context is different? You don't need answers in week 1, but you need to be thinking about this from the start, because it shapes what you build and how you build it.

Kickstart Mixpanel Experiments

Mixpanel Experiments is how we'll run every A/B test going forward. It's built into Mixpanel — Feature Flags control who sees what, exposure events are tracked automatically, and the Experiment report measures lift and significance against any metric we already track. The infrastructure is there. What's missing is someone who digs in, understands how it works, and starts building toward our first experiment.

Your week 1 job is to become the person on the team who knows this tool inside out. Get into our Mixpanel project, understand the events we already track and how the activation funnel is structured. Then go deep on how Mixpanel Experiments works — how Feature Flags are created, how the SDK assigns users to variants, how exposure events flow into the Experiment report, what the difference is between Sequential and Frequentist models, and when to use which.

By the end of the week, come back with a clear picture of what we need to do to run our first experiment. What's already in place, what needs engineering work in the SDK, what decisions we need to make together. That assessment is the deliverable — not a finished implementation, but a concrete plan that lets us move fast in week 2.

Start here:

- Our Mixpanel project (ID `4011597`) — explore the events, the activation funnel, and the custom events. And an important note here. We just migrated from our old project since those included Auth0 IDs, which we didn't want to have in Mixpanel.
- Mixpanel Experiments docs: Experiments: Measure the impact of a/b testing - Mixpanel Docs
- Feature Flags overview: How to Safely Ship New Features with Flags - Mixpanel Docs
- Experiment design best practices: Drive Product Innovation with Mixpanel Experiments - Mixpanel Docs

---

Day 1

By the end of today, you have a working laptop, access to all your tools, you understand what Ditto does and how your domain works, and you know your mission for the week.

- [ ] Laptop setup: everything works
- [ ] Company-wide tool access verified
- [ ] Domain-specific tool access verified
- [ ] Set your email signature; ask a colleague to send you an email, then copy it from there
- [ ] Founder talk: why Ditto exists, where we're going
- [ ] Domain introduction with lead: your team, OKR, how we work
- [ ] Walk through this checklist with the lead: fill in anything specific to you still missing
- [ ] Governance sign-off with ops @Bart Voorn
- [ ] Understand your **week-1 key result** - what you're delivering and demo-ing on Friday
- [ ] Prep for tomorrow's standup: what did you do today, what's the plan for tomorrow?

Week 1

No fixed order, work through these during the week.

Get to know Ditto

- [ ] Coffee chat with every team member, you schedule these yourself
- [ ] Security onboarding session
- [ ] Read through company-wide documentation. You can check everything you want, but these are a priority:
  - [ ] Ways of Working
  - [ ] Old office 🏢 (Microlab)
  - [ ] OKRs at Ditto
  - [ ] Focus on Life Moments in Healthcare
  - [ ] Security Guidelines for All
  - [ ] Product Strategy
  - [ ] Data Breach Management Policy

- [ ] Understand how your domain does standups, retros, and planning
- [ ] Read through domain-specific documentation:
  - [ ] [domain fills in: e.g., architecture overview, coding guidelines, design system, ...]
- [ ] Get to know your Growth Agents (Pokemons) and Cursor Automations
- [ ] Walkthrough of Customer.io
- [ ] Walkthrough of Mixpanel
- [ ] Understand how Growth operates agentically with Mew
  - [ ] Niek needs to figure out how to make his locally hosted second brain accessible.

Get to know your role

- [ ] **"Who are you" conversation** with lead (day 3 or 4):
  - [ ] What are you good at? This doesn't have to be within your role at Ditto, if you see a way to add value to another domain, you're encouraged to do so.
- [ ] Where do you want to grow / learn / develop?
- [ ] What is the most productive way for you to work? This becomes a reference for you and your team.

> ℹ️ We will create a document out of that, and the same document will be used to “learn” together. See where you thrive, where you want to learn, and where you may need to learn.

Your mission

- Work on your **week-1 key result**
- Ask questions. Nobody expects you to figure everything out alone
- **Friday: demo your key result + prep retro**

Wrap up

- [ ] Update your LinkedIn profile to reflect your role at Ditto
- [ ] Share a post with Niek about your hire. Make it personal. We will increase LinkedIn content and aim for every employee to become a thought leader.

First Month

- Continue working on your OKR
- **3-week evaluation conversation**, A check-in with 3 people (lead, founder, one cross-functional peer). Topics:
  - What has gone well?
  - What has been challenging?
  - Are you on track? Any concerns from either side?
  - Goals and focus areas moving forward

First-month milestones

By the end of your first month, three things should be true.

1. **Mixpanel Experiments is ready to run — and your first experiments are shipping.** You've gone from reading the docs to having the experimentation engine moving. You understand how Feature Flags, exposure events, and the Experiment report work together. You've identified what needs engineering work in the SDK, and you're shipping your first activation experiments. This connects directly to the Q2 OKRs: deliverable 14 under O2-KR2 calls for you to run at least 2 activation experiments by the end of May, targeting the biggest drop-off points in the onboarding funnel. That deadline is real — it gives us 4+ weeks of data before the Q2 close, allowing us to measure impact and iterate.
2. **The first version of a growth model is taking shape.** Together with Niek, you've started mapping out which user behaviors predict retention — beyond the summary insight we already have. What role does the care circle play? What distinguishes users who retain from those who don't? What are the candidate levers we could test? This isn't a finished model — it's the first draft, built from the data we have and the questions we can't answer yet.
3. **You own the data quality standard for product releases.** Every feature that ships at Ditto should be trackable and measurable. You're the person who makes sure that's true. That means reviewing tracking plans before engineering builds, making sure events and properties follow our naming conventions, and catching gaps before they reach production. This is foundational — because every experiment you design starts with a hypothesis, and every hypothesis needs a measurement plan. If the data isn't clean, nothing downstream works.

Brainpicks and writing

There's a layer to this role that's different from most engineering jobs. Beyond Mixpanel, Jira, and the usual tools, we run a system we're actively building as we go — and part of your job is to help decide what that system becomes.

The first layer is Mew — a workspace in Cursor that acts as the operational brain for growth. It's where the AI agents live, where skills and automations are defined, where the experiment pipeline is orchestrated, and where institutional knowledge accumulates. Six agents run autonomously in Slack, each owning a domain: analytics, experimentation, retention, feedback, expansion, and product quality. They coordinate with each other, surface insights, and flag problems. You'll work with them daily, but more importantly, you'll help shape how they evolve.

The second layer is the growth second brain — currently an Obsidian vault that Niek maintains using the LLM Wiki pattern. Raw sources go in (meeting notes from Granola, experiment results, articles, strategic decisions), and the system structures them into queryable wiki pages with cross-references. Right now it's single-player. With your joining, it becomes shared.

The vision is that everything we learn compounds. Every meeting we have gets captured and ingested. Every experiment — hypothesis, result, verdict — becomes a queryable source. Every article we clip, every Mixpanel deep-dive, every user interview gets structured and connected. Over time, you can ask "what do we know about caretaker activation?" and get a synthesized answer drawn from months of compounding evidence, not just what one of us happens to remember.

We're essentially building our own AI-native growth infrastructure — detection, reasoning, and action connected through a knowledge layer. Obsidian holds the knowledge, Cursor orchestrates, and the agent network executes. It's early, it's opinionated, and it's one of the more unusual things about working here.

Things to chew on together:

About Mew and the agent network:

- Six agents running autonomously in Slack — is that the right architecture, or should some of them be merged, split, or rebuilt? Where do you see overlap or gaps?
- The agents are prompt-driven (markdown files in `automations/`). What happens when we want them to learn from their own outputs? Right now they don't — every run starts fresh. Should they?
- Cursor is the orchestration layer, but it's designed for code. Are we stretching it beyond what it's good at, or is "everything is a markdown workspace" actually the right abstraction for a growth team?
- The two-tier model (Level 1: autonomous in Slack, Level 2: interactive in Cursor) — does that split make sense, or should there be a third tier? What would a fully autonomous agent that runs experiments end-to-end even look like?

About the second brain:

- The LLM Wiki pattern works for structured ingestion, but it's lossy — the AI decides what matters when it summarizes a source. How do we make sure we're not throwing away the insights that matter most?
- Right now knowledge flows in one direction: raw → wiki → query. What would it look like if the wiki could trigger actions — like flagging when a new experiment result contradicts an existing belief?

- How do we make this genuinely shared without it becoming noisy? Two people contributing is manageable. What if the team grows to five?

About build versus buy:

- Hilbert offers an AI-native growth infrastructure — anomaly detection, root cause analysis, predictive LTV, churn scoring, automated retention plays. Their case studies are Blank Street, FreshDirect, Erewhon. Could their ML models work at our scale (1,739 Core MAU), or do we need 10x more data before predictions are useful?
- Edra reverse-engineers how your team operates by reading tickets, logs, and messages, then turns that into executable playbooks for AI agents. That's close to what the second brain does, but automated. Could it replace or accelerate what we're building manually? Or is the value precisely that we control the structure?
- Where's the line between building growth infrastructure as a competitive advantage versus overinvesting in tooling when we should be running experiments?
- If we had a budget for one external tool to plug into our stack, where would it create the most leverage — detection, knowledge management, or execution?

About the bigger picture:

- We're a two-person growth team with an agent network that would normally require a team of six. What breaks first as we scale — the agents, the knowledge base, or the humans?
- Most growth teams are organized around functions (acquisition, retention, analytics). We're organized around a shared system. Is that better, or are we creating a single point of failure?
- At what point does this stop being "Niek's workspace that Wouter uses" and become a proper platform? What would that transition actually require?

This is where your AI engineering background becomes a multiplier. You're not just a user of this system — you can evaluate it technically, challenge assumptions, and help decide where we build versus buy. Niek will onboard you into Mew and the second brain step by step. The system is built to grow with the team, and you're joining at the point where it stops being one person's setup and starts becoming something bigger.